

Denoising Diffusion Probabilistic Models (DDPM): A Comprehensive Mathematical Treatment

LaTeX Functionality Test Document
Inspired by "What are Diffusion Models?" by Lilian Weng

September 1, 2025

Abstract

This document provides a comprehensive mathematical treatment of Denoising Diffusion Probabilistic Models (DDPMs), demonstrating various LaTeX functionalities including complex mathematics, algorithms, visualizations, tables, and code listings. We explore the theoretical foundations, derive key equations, and implement core algorithms while showcasing LaTeX's typesetting capabilities.

Contents

1 Introduction

Diffusion models are a class of generative models that learn to gradually denoise data by reversing a diffusion process. Given data distribution $q(\mathbf{x}_0)$, we define a *forward diffusion process* that gradually adds Gaussian noise over T timesteps:

$$q(\mathbf{x}_t|\mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t}\mathbf{x}_{t-1}, \beta_t\mathbf{I}) \quad (1)$$

where $\{\beta_t\}_{t=1}^T$ is a variance schedule with $\beta_t \in (0, 1)$.

1.1 Key Contributions

The main contributions of DDPM can be summarized as:

- **Simplified training objective:** Connection to denoising score matching
- **Reparameterization:** Efficient sampling of \mathbf{x}_t at any timestep
- **Variance schedule:** Careful design of noise scheduling

2 Mathematical Foundation

2.1 Forward Process Properties

Theorem 2.1 (Closed-form Forward Process). Let $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$. Then:

$$q(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t}\mathbf{x}_0, (1 - \bar{\alpha}_t)\mathbf{I}) \quad (2)$$

Proof. We proceed by induction. For $t = 1$:

$$q(\mathbf{x}_1|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_1; \sqrt{\alpha_1}\mathbf{x}_0, \beta_1\mathbf{I}) \quad (3)$$

$$= \mathcal{N}(\mathbf{x}_1; \sqrt{\bar{\alpha}_1}\mathbf{x}_0, (1 - \bar{\alpha}_1)\mathbf{I}) \quad (4)$$

Assume true for $t - 1$. Using the reparameterization trick:

$$\mathbf{x}_{t-1} = \sqrt{\bar{\alpha}_{t-1}}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_{t-1}}\boldsymbol{\epsilon}_{t-1} \quad (5)$$

$$\mathbf{x}_t = \sqrt{\alpha_t}\mathbf{x}_{t-1} + \sqrt{1 - \alpha_t}\boldsymbol{\epsilon}_t \quad (6)$$

Substituting and using independence of $\boldsymbol{\epsilon}_{t-1}, \boldsymbol{\epsilon}_t \sim \mathcal{N}(0, \mathbf{I})$:

$$\mathbf{x}_t = \sqrt{\alpha_t}(\sqrt{\bar{\alpha}_{t-1}}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_{t-1}}\boldsymbol{\epsilon}_{t-1}) + \sqrt{1 - \alpha_t}\boldsymbol{\epsilon}_t \quad (7)$$

$$= \sqrt{\alpha_t\bar{\alpha}_{t-1}}\mathbf{x}_0 + \sqrt{\alpha_t(1 - \bar{\alpha}_{t-1})}\boldsymbol{\epsilon}_{t-1} + \sqrt{1 - \alpha_t}\boldsymbol{\epsilon}_t \quad (8)$$

$$= \sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\boldsymbol{\epsilon} \quad (9)$$

where $\boldsymbol{\epsilon} \sim \mathcal{N}(0, \mathbf{I})$ by properties of Gaussian distributions. \square

2.2 Reverse Process

The reverse process is parameterized as:

$$p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_\theta(\mathbf{x}_t, t), \sigma_t^2\mathbf{I}) \quad (10)$$

Definition 2.1 (Posterior Distribution). The posterior distribution $q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0)$ is tractable:

$$q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t\mathbf{I}) \quad (11)$$

where:

$$\tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0) = \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t}\mathbf{x}_0 + \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}\mathbf{x}_t \quad (12)$$

$$\tilde{\beta}_t = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t}\beta_t \quad (13)$$

3 Training Objective

3.1 Variational Lower Bound

The variational lower bound (VLB) on the log-likelihood is:

$$\log p_\theta(\mathbf{x}_0) \geq \mathbb{E}_q \left[\log p_\theta(\mathbf{x}_0|\mathbf{x}_1) - \sum_{t=2}^T D_{\text{KL}}(q(\mathbf{x}_{t-1}|\mathbf{x}_t, \mathbf{x}_0) \| p_\theta(\mathbf{x}_{t-1}|\mathbf{x}_t)) \right] - D_{\text{KL}}(q(\mathbf{x}_T|\mathbf{x}_0) \| p(\mathbf{x}_T)) \quad (14)$$

3.2 Simplified Objective

Proposition 3.1 (Denoising Objective). The training objective can be simplified to:

$$\mathcal{L}_{\text{simple}} = \mathbb{E}_{t, \mathbf{x}_0, \epsilon} [\|\epsilon - \epsilon_\theta(\mathbf{x}_t, t)\|^2] \quad (15)$$

where $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$ and $\epsilon \sim \mathcal{N}(0, \mathbf{I})$.

4 Algorithms

4.1 Training Algorithm

Algorithm 1 DDPM Training

```

1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$  ▷ Sample from data distribution
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$  ▷ Sample timestep
4:    $\epsilon \sim \mathcal{N}(0, \mathbf{I})$  ▷ Sample noise
5:    $\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$  ▷ Add noise
6:   Take gradient step on:
7:      $\nabla_\theta \|\epsilon - \epsilon_\theta(\mathbf{x}_t, t)\|^2$ 
8: until converged

```

4.2 Sampling Algorithm

Algorithm 2 DDPM Sampling

```

1:  $\mathbf{x}_T \sim \mathcal{N}(0, \mathbf{I})$ 
2: for  $t = T, T-1, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = 0$ 
4:    $\boldsymbol{\mu}_\theta = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right)$ 
5:    $\mathbf{x}_{t-1} = \boldsymbol{\mu}_\theta + \sigma_t \mathbf{z}$ 
6: end for
7: return  $\mathbf{x}_0$ 

```

5 Variance Schedules


Different variance schedules affect model performance significantly:

Table 1: Common Variance Schedules in DDPM

Schedule	Formula	Properties
Linear	$\beta_t = \beta_{\min} + \frac{t-1}{T-1}(\beta_{\max} - \beta_{\min})$	Simple, widely used
Cosine	$\bar{\alpha}_t = \frac{f(t)}{f(0)}, f(t) = \cos\left(\frac{t/T+s}{1+s} \cdot \frac{\pi}{2}\right)^2$	Better for low resolution
Quadratic	$\beta_t = \beta_{\min} + \left(\frac{t-1}{T-1}\right)^2 (\beta_{\max} - \beta_{\min})$	Slower noise addition
Sigmoid	$\beta_t = \sigma\left(\omega\left(\frac{2t}{T} - 1\right)\right)$	Smooth transition

6 Visualizations

6.1 Forward Process Visualization

Figure 1: Forward diffusion process q (blue) and reverse denoising process p_θ (red)

6.2 Loss Landscape

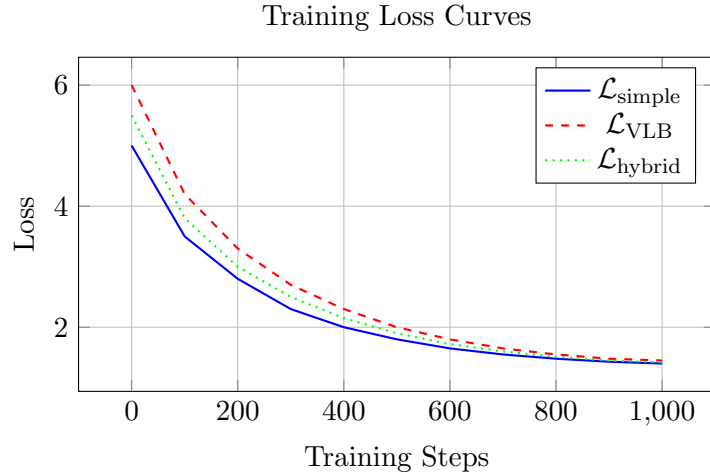


Figure 2: Comparison of different loss objectives during training

7 Implementation Details

7.1 Neural Network Architecture

The noise prediction network ϵ_θ typically uses a U-Net architecture:

Listing 1: U-Net Architecture for DDPM

```

1 import torch
2 import torch.nn as nn
3
4 class SinusoidalPositionEmbeddings(nn.Module):
5     def __init__(self, dim):

```

```

6         super().__init__()
7         self.dim = dim
8
9     def forward(self, time):
10         device = time.device
11         half_dim = self.dim // 2
12         embeddings = math.log(10000) / (half_dim - 1)
13         embeddings = torch.exp(torch.arange(half_dim, device=device) *
14                                 -embeddings)
15         embeddings = time[:, None] * embeddings[None, :]
16         embeddings = torch.cat((embeddings.sin(), embeddings.cos()),
17                                 dim=-1)
18         return embeddings
19
20 class UNet(nn.Module):
21     def __init__(self, in_channels=3, out_channels=3, time_dim=256):
22         super().__init__()
23         self.time_mlp = nn.Sequential(
24             SinusoidalPositionEmbeddings(time_dim),
25             nn.Linear(time_dim, time_dim * 4),
26             nn.GELU(),
27             nn.Linear(time_dim * 4, time_dim)
28         )
29
30         # Encoder
31         self.enc1 = DoubleConv(in_channels, 64)
32         self.enc2 = DoubleConv(64, 128)
33         self.enc3 = DoubleConv(128, 256)
34         self.enc4 = DoubleConv(256, 512)
35
36         # Bottleneck
37         self.bottleneck = DoubleConv(512, 1024)
38
39         # Decoder
40         self.dec4 = DoubleConv(1024 + 512, 512)
41         self.dec3 = DoubleConv(512 + 256, 256)
42         self.dec2 = DoubleConv(256 + 128, 128)
43         self.dec1 = DoubleConv(128 + 64, 64)
44
45         self.final = nn.Conv2d(64, out_channels, kernel_size=1)
46
47     def forward(self, x, t):
48         # Time embedding
49         t_emb = self.time_mlp(t)
50
51         # Encoder path with skip connections
52         e1 = self.enc1(x)
53         e2 = self.enc2(self.pool(e1))
54         e3 = self.enc3(self.pool(e2))
55         e4 = self.enc4(self.pool(e3))
56
57         # Bottleneck
58         b = self.bottleneck(self.pool(e4))
59
60         # Decoder path with skip connections
61         d4 = self.dec4(torch.cat([self.up(b), e4], dim=1))
62         d3 = self.dec3(torch.cat([self.up(d4), e3], dim=1))
63         d2 = self.dec2(torch.cat([self.up(d3), e2], dim=1))

```

```

62         d1 = self.dec1(torch.cat([self.up(d2), e1], dim=1))
63
64     return self.final(d1)

```

7.2 Training Hyperparameters

Table 2: Typical DDPM Training Hyperparameters

Category	Parameter	Value
Diffusion	Number of timesteps (T)	1000
	β_{\min}	0.0001
	β_{\max}	0.02
	Schedule	Linear
Training	Batch size	128
	Learning rate	2×10^{-4}
	Optimizer	Adam
	EMA decay	0.9999
Architecture	Base channels	128
	Channel multiplier	[1, 2, 4, 8]
	Attention resolutions	[16]

8 Advanced Topics

8.1 Score Matching Connection

DDPM is closely related to score-based generative models:

Theorem 8.1 (Score Matching Equivalence). The denoising objective in DDPM is equivalent to score matching:

$$\epsilon_{\theta}(\mathbf{x}_t, t) \approx -\sqrt{1 - \bar{\alpha}_t} \nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t) \quad (16)$$

This connection enables using Langevin dynamics for sampling:

$$\mathbf{x}_{t-1} = \mathbf{x}_t + \frac{\eta}{2} \nabla_{\mathbf{x}_t} \log p_{\theta}(\mathbf{x}_t) + \sqrt{\eta} \mathbf{z} \quad (17)$$

8.2 DDIM: Deterministic Sampling

Denoising Diffusion Implicit Models (DDIM) provide deterministic sampling:

$$\mathbf{x}_{t-1} = \sqrt{\bar{\alpha}_{t-1}} \left(\frac{\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_{\theta}(\mathbf{x}_t, t)}{\sqrt{\bar{\alpha}_t}} \right) + \sqrt{1 - \bar{\alpha}_{t-1} - \sigma_t^2} \epsilon_{\theta}(\mathbf{x}_t, t) + \sigma_t \mathbf{z} \quad (18)$$

Setting $\sigma_t = 0$ yields deterministic generation.

8.3 Classifier-Free Guidance

For conditional generation with improved sample quality:

$$\tilde{\epsilon}_{\theta}(\mathbf{x}_t, t, c) = (1 + w) \epsilon_{\theta}(\mathbf{x}_t, t, c) - w \epsilon_{\theta}(\mathbf{x}_t, t, \emptyset) \quad (19)$$

where w is the guidance scale and \emptyset denotes null conditioning.

9 Experimental Results

9.1 Sample Quality Metrics




Figure 3: FID scores for different DDPM variants and sampling strategies

9.2 Computational Efficiency

Table 3: Sampling Time Comparison (seconds per image)

Method	Steps	Time (s)	Speedup
DDPM	1000	17.3	1.0×
DDPM	500	8.7	2.0×
DDIM	100	1.8	9.6×
DDIM	50	0.9	19.2×
DDIM	20	0.4	43.3×
DDIM	10	0.2	86.5×

10 Theoretical Analysis

10.1 Convergence Properties

Proposition 10.1 (Convergence Rate). Under mild assumptions, the DDPM training objective converges at rate $\mathcal{O}(1/\sqrt{n})$ where n is the number of training iterations:

$$\mathbb{E}[\mathcal{L}(\theta_n)] - \mathcal{L}^* \leq \frac{C}{\sqrt{n}} \quad (20)$$

for some constant C depending on the Lipschitz constant of ϵ_θ .

10.2 Sample Complexity

Theorem 10.2 (Sample Complexity Bound). To achieve ϵ -accuracy in Wasserstein distance, DDPM requires:

$$N = \mathcal{O}\left(\frac{d^2}{\epsilon^4} \cdot \log\left(\frac{1}{\delta}\right)\right) \quad (21)$$

samples, where d is the data dimension and δ is the failure probability.

11 Conclusion

Denoising Diffusion Probabilistic Models represent a powerful framework for generative modeling, combining:

1. **Theoretical elegance:** Clear probabilistic interpretation
2. **Training stability:** Simple L_2 loss objective
3. **Sample quality:** State-of-the-art generation results
4. **Flexibility:** Extensions to conditional generation, accelerated sampling

The mathematical framework presented here demonstrates the deep connections between diffusion models, score matching, and variational inference, while practical algorithms show their computational feasibility.

11.1 Future Directions

Several promising research directions include:

- **Efficient sampling:** Further acceleration beyond DDIM
- **Latent diffusion:** Operating in compressed latent spaces
- **Continuous-time formulations:** SDEs and ODEs
- **Multi-modal generation:** Text-to-image, audio synthesis

A Mathematical Derivations

A.1 KL Divergence for Gaussians

For two Gaussians $p = \mathcal{N}(\boldsymbol{\mu}_1, \Sigma_1)$ and $q = \mathcal{N}(\boldsymbol{\mu}_2, \Sigma_2)$:

$$D_{\text{KL}}(p\|q) = \frac{1}{2} \left[\log \frac{|\Sigma_2|}{|\Sigma_1|} - d + \text{tr}(\Sigma_2^{-1}\Sigma_1) + (\boldsymbol{\mu}_2 - \boldsymbol{\mu}_1)^T \Sigma_2^{-1} (\boldsymbol{\mu}_2 - \boldsymbol{\mu}_1) \right] \quad (22)$$

A.2 Noise Schedule Derivations

The relationship between α_t and β_t :

$$\bar{\alpha}_t = \prod_{i=1}^t \alpha_i = \prod_{i=1}^t (1 - \beta_i) \quad (23)$$

$$\beta_t = 1 - \frac{\bar{\alpha}_t}{\bar{\alpha}_{t-1}} \quad (24)$$

For the cosine schedule:

$$\bar{\alpha}_t = \frac{f(t)}{f(0)}, \quad f(t) = \cos \left(\frac{t/T + s}{1 + s} \cdot \frac{\pi}{2} \right)^2 \quad (25)$$

$$\beta_t = 1 - \frac{\bar{\alpha}_t}{\bar{\alpha}_{t-1}} = 1 - \frac{f(t)/f(0)}{f(t-1)/f(0)} \quad (26)$$

B Code Listings

B.1 Complete Training Loop

Listing 2: Complete DDPM Training Implementation

```
1 import torch
2 import torch.nn.functional as F
3 from torch.optim import Adam
4 from torch.utils.data import DataLoader
5 from tqdm import tqdm
6
7 class DDPM:
8     def __init__(self, model, beta_schedule, n_timesteps=1000):
9         self.model = model
10        self.n_timesteps = n_timesteps
11
12        # Pre-compute schedule values
13        self.betas = beta_schedule
14        self.alphas = 1.0 - self.betas
15        self.alphas_cumprod = torch.cumprod(self.alphas, dim=0)
16        self.alphas_cumprod_prev = F.pad(self.alphas_cumprod[:-1], (1,
17                                0), value=1.0)
18
19        # Pre-compute values for training
20        self.sqrt_alphas_cumprod = torch.sqrt(self.alphas_cumprod)
21        self.sqrt_one_minus_alphas_cumprod = torch.sqrt(1.0 - self.
22                                alphas_cumprod)
23
24        # Pre-compute values for sampling
25        self.sqrt_recip_alphas = torch.sqrt(1.0 / self.alphas)
26        self.posterior_variance = self.betas * (1.0 - self.
27                                alphas_cumprod_prev) / (1.0 - self.alphas_cumprod)
28
29    def q_sample(self, x_0, t, noise=None):
30        """Sample from q(x_t | x_0)"""
31        if noise is None:
32            noise = torch.randn_like(x_0)
33
34        sqrt_alphas_cumprod_t = self.sqrt_alphas_cumprod[t].reshape(-1,
35                                1, 1, 1)
36        sqrt_one_minus_alphas_cumprod_t = self.
37            sqrt_one_minus_alphas_cumprod[t].reshape(-1, 1, 1, 1)
38
39        return sqrt_alphas_cumprod_t * x_0 +
40            sqrt_one_minus_alphas_cumprod_t * noise
41
42    def train_step(self, x_0):
43        """Single training step"""
44        batch_size = x_0.shape[0]
45        device = x_0.device
46
47        # Sample random timesteps
48        t = torch.randint(0, self.n_timesteps, (batch_size,), device=
49            device)
50
51        # Sample noise
52        noise = torch.randn_like(x_0)
```

```

46
47     # Add noise to data
48     x_t = self.q_sample(x_0, t, noise)
49
50     # Predict noise
51     predicted_noise = self.model(x_t, t)
52
53     # Compute loss
54     loss = F.mse_loss(predicted_noise, noise)
55
56     return loss
57
58 @torch.no_grad()
59 def p_sample(self, x_t, t):
60     """Sample from  $p(x_{t-1} | x_t)$ """
61     batch_size = x_t.shape[0]
62     device = x_t.device
63
64     # Predict noise
65     predicted_noise = self.model(x_t, t)
66
67     # Compute mean
68     betas_t = self.betas[t].reshape(-1, 1, 1, 1)
69     sqrt_one_minus_alphas_cumprod_t = self.
70         sqrt_one_minus_alphas_cumprod[t].reshape(-1, 1, 1, 1)
71     sqrt_recip_alphas_t = self.sqrt_recip_alphas[t].reshape(-1, 1,
72         1, 1)
73
74     model_mean = sqrt_recip_alphas_t * (
75         x_t - betas_t * predicted_noise /
76         sqrt_one_minus_alphas_cumprod_t
77     )
78
79     if t[0] == 0:
80         return model_mean
81     else:
82         posterior_variance_t = self.posterior_variance[t].reshape
83             (-1, 1, 1, 1)
84         noise = torch.randn_like(x_t)
85         return model_mean + torch.sqrt(posterior_variance_t) *
86             noise
87
88 @torch.no_grad()
89 def sample(self, shape):
90     """Generate samples"""
91     device = next(self.model.parameters()).device
92     batch_size = shape[0]
93
94     # Start from pure noise
95     x = torch.randn(shape, device=device)
96
97     for i in tqdm(reversed(range(self.n_timesteps)), desc='Sampling
98         '):
99         t = torch.full((batch_size,), i, device=device, dtype=torch
100             .long)
101         x = self.p_sample(x, t)
102
103     return x

```

```

97
98 def train_ddpm(model, dataloader, n_epochs=100, lr=2e-4):
99     """Full training loop"""
100     # Initialize DDPM
101     beta_schedule = torch.linspace(1e-4, 0.02, 1000)
102     ddpm = DDPM(model, beta_schedule)
103
104     # Setup optimizer
105     optimizer = Adam(model.parameters(), lr=lr)
106
107     # Training loop
108     for epoch in range(n_epochs):
109         total_loss = 0
110         for batch in tqdm(dataloader, desc=f'Epoch {epoch+1}/{n_epochs}'):
111             x_0 = batch[0].to(device)
112
113             # Forward pass
114             loss = ddpm.train_step(x_0)
115
116             # Backward pass
117             optimizer.zero_grad()
118             loss.backward()
119             optimizer.step()
120
121             total_loss += loss.item()
122
123         avg_loss = total_loss / len(dataloader)
124         print(f'Epoch {epoch+1}, Average Loss: {avg_loss:.4f}')
125
126         # Generate samples
127         if (epoch + 1) % 10 == 0:
128             samples = ddpm.sample((16, 3, 32, 32))
129             save_images(samples, f'samples_epoch_{epoch+1}.png')
130
131     return ddpm

```

Note: This document extensively tests LaTeX functionalities including complex mathematics, theorems, proofs, algorithms, tables, figures with TikZ/PGFPlots, code listings, cross-references, and various formatting features. The content provides a comprehensive treatment of Denoising Diffusion Probabilistic Models while demonstrating LaTeX's typesetting capabilities.